

# AML Rohit Assignment-4

ROHIT VURADI

2024-04-21

## Introduction

This demonstrates how to perform sentiment analysis on the IMDB movie review dataset using Recurrent Neural Networks (RNNs) with the Keras library in R. We'll modify the example from Chapter 6 of the book to align with the specific requirements of our assignment.

RNNs and Transformers are the commonly used models in the field of NLP ( Natural Language Processing). They are used for processing sequential data such as text to derive meaning/understanding. Else GRUs, including Long Short-Term Memory (LSTM) and Gated Recurrent Unit (GRU), are very capable of keeping a memory over time since they preserve the memory by sequencing. And different, to transformation which apply attention self-mechanisms to the model work, there is a capturing of long-range aspects in sequences. RNNs and transformers can be used when it comes to text and data in sequences both for tasks such as sentiment analysis, machine translation, and text generation.

### To apply RNNs or Transformers to text data: To apply RNNs or Transformers to text data:

Preprocess the text data: Let's take the text and break it down into words or subwords, convert those tokens into numerical form (such as word embeddings) , and finally ensure the sequences of the sequence length are equal.

Define the model architecture: Construct a neural network model, which is going to have suitable layers for processing sequential data. For RNNs this is where it embeds the words in different ways and then these are concatenated (LSTM, GRU). To the Transformers, forming a complementary relation between the self-attention layers and the feedforward layers would be desirable.

Compile and train the model: Collect the model with the proper loss function and optimizer, then provide the training data for the model successively.

Evaluate the model: Evaluate the model on the validation or test set, which is different from the training set and contains the information about how well the model performs at text data processing.

```
# load the neccessary library  
library(keras)
```

## Loading the IMDB Dataset

We'll load the IMDB dataset with the following modifications:

Cutoff reviews after 150 words. Restrict training samples to 100. Validation on 10,000 samples. Consider only the top 10,000 words.

```
# load the IMDB dataset  
imdb <- dataset_imdb(num_words = 10000)  
maxlen <- 150  
x_train <- pad_sequences(imdb$train$x, maxlen = maxlen)  
x_test <- pad_sequences(imdb$test$x, maxlen = maxlen)  
sample_size <- 100
```

```

x_train <- x_train[1:sample_size, ]
y_train <- as.matrix(imdb$train$y)[1:sample_size, ]
x_val <- x_test[1:10000, ]
y_val <- as.matrix(imdb$test$y)[1:10000, ]

# check the dimensions of the training and validation sets
dim(x_train)

## [1] 100 150

dim(y_train)

## NULL

dim(x_val)

## [1] 10000 150

dim(y_val)

## NULL

# check the first 5 elements of the training set
head(x_train)


##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12] [,13] [,14]
## [1,]  12  16  43 530  38  76  15  13 1247    4  22  17  515  17
## [2,]   4 249 126  93   4 114   9 2300 1523    5 647   4  116   9
## [3,]   0   0   0   0   0   0   0   0   0    1  14  47   8  30
## [4,]   4   2   7 154   5   4 518  53   2    2   7 3211 882  11
## [5,]   0   0   0   1 249 1323   7  61 113   10  10  13 1637  14
## [6,]   0   0   0   0   0   0   0   0   0    0   0   0   0   0
##      [,15] [,16] [,17] [,18] [,19] [,20] [,21] [,22] [,23] [,24] [,25] [,26]
## [1,]  12  16  626  18   2   5  62  386  12   8  316   8
## [2,]  35 8163   4  229   9  340 1322   4  118   9   4  130
## [3,]  31   7   4  249 108   7   4 5974  54  61  369  13
## [4,] 399  38  75  257 3807  19   2  17  29 456   4  65
## [5,]  20  56  33 2401  18 457  88  13 2626 1400  45 3171
## [6,]   0   0   0   0   0   0   0   0   0   0   0   0
##      [,27] [,28] [,29] [,30] [,31] [,32] [,33] [,34] [,35] [,36] [,37] [,38]
## [1,]  106   5   4 2223 5244  16  480  66 3785  33   4  130
## [2,] 4901  19   4 1002   5  89  29 952  46  37   4 455
## [3,]  71 149  14  22  112   4 2401 311  12  16 3711  33
## [4,]   7  27 205 113  10  10   2   4   2   2   9 242
## [5,]  13  70  79  49 706 919  13  16 355 340 355 1696
## [6,]   0   0   0   0   0   0   0   0   0   0   0   0
##      [,39] [,40] [,41] [,42] [,43] [,44] [,45] [,46] [,47] [,48] [,49] [,50]
## [1,]  12  16  38  619   5  25 124  51  36 135  48  25
## [2,]   9  45  43  38 1543 1905 398   4 1649  26 6853   5
## [3,]  75  43 1829 296   4  86 320  35 534  19 263 4821
## [4,]   4  91 1202   2   5 2070 307  22   7 5168 126  93
## [5,]  96 143   4  22  32 289   7  61 369  71 2359   5
## [6,]   0   0   0   0   0   0   0   0   0   0   0   0
##      [,51] [,52] [,53] [,54] [,55] [,56] [,57] [,58] [,59] [,60] [,61] [,62]
## [1,] 1415  33   6  22  12 215  28  77  52   5  14 407
## [2,] 163  11 3215   2   4 1153   9 194 775   7 8255   2
## [3,] 1301   4 1873  33  89  78  12  66  16   4  360   7

```

##	[4,]	40	2	13	188	1076	3222	19	4	2	7	2348	537
##	[5,]	13	16	131	2073	249	114	249	229	249	20	13	28
##	[6,]	0	0	0	0	0	0	0	0	0	0	0	0
##		[,63]	[,64]	[,65]	[,66]	[,67]	[,68]	[,69]	[,70]	[,71]	[,72]	[,73]	[,74]
##	[1,]	16	82	2	8	4	107	117	5952	15	256	4	2
##	[2,]	349	2637	148	605	2	8003	15	123	125	68	2	6853
##	[3,]	4	58	316	334	11	4	1716	43	645	662	8	257
##	[4,]	23	53	537	21	82	40	2	13	2	14	280	13
##	[5,]	126	110	13	473	8	569	61	419	56	429	6	1513
##	[6,]	0	0	0	0	0	0	0	0	0	0	0	0
##		[,75]	[,76]	[,77]	[,78]	[,79]	[,80]	[,81]	[,82]	[,83]	[,84]	[,85]	[,86]
##	[1,]	7	3766	5	723	36	71	43	530	476	26	400	317
##	[2,]	15	349	165	4362	98	5	4	228	9	43	2	1157
##	[3,]	85	1200	42	1228	2578	83	68	3912	15	36	165	1539
##	[4,]	219	4	2	431	758	859	4	953	1052	2	7	5991
##	[5,]	18	35	534	95	474	570	5	25	124	138	88	12
##	[6,]	0	0	0	0	0	0	0	0	0	0	0	0
##		[,87]	[,88]	[,89]	[,90]	[,91]	[,92]	[,93]	[,94]	[,95]	[,96]	[,97]	[,98]
##	[1,]	46	7	4	2	1029	13	104	88	4	381	15	297
##	[2,]	15	299	120	5	120	174	11	220	175	136	50	9
##	[3,]	278	36	69	2	780	8	106	14	6905	1338	18	6
##	[4,]	5	94	40	25	238	60	2	4	2	804	2	7
##	[5,]	421	1543	52	725	6397	61	419	11	13	1571	15	1543
##	[6,]	0	0	0	0	0	0	0	0	0	0	0	0
##		[,99]	[,100]	[,101]	[,102]	[,103]	[,104]	[,105]	[,106]	[,107]	[,108]		
##	[1,]	98	32	2071	56	26	141	6	194	7486	18		
##	[2,]	4373	228	8255	5	2	656	245	2350	5	4		
##	[3,]	22	12	215	28	610	40	6	87	326	23		
##	[4,]	4	9941	132	8	67	6	22	15	9	283		
##	[5,]	20	11	4	2	5	296	12	3524	5	15		
##	[6,]	0	0	0	0	0	0	0	0	0	1		
##		[,109]	[,110]	[,111]	[,112]	[,113]	[,114]	[,115]	[,116]	[,117]	[,118]		
##	[1,]	4	226	22	21	134	476	26	480	5	144		
##	[2,]	9837	131	152	491	18	2	32	7464	1212	14		
##	[3,]	2300	21	23	22	12	272	40	57	31	11		
##	[4,]	8	5168	14	31	9	242	955	48	25	279		
##	[5,]	421	128	74	233	334	207	126	224	12	562		
##	[6,]	778	128	74	12	630	163	15	4	1766	7982		
##		[,119]	[,120]	[,121]	[,122]	[,123]	[,124]	[,125]	[,126]	[,127]	[,128]		
##	[1,]	30	5535	18	51	36	28	224	92	25	104		
##	[2,]	9	6	371	78	22	625	64	1382	9	8		
##	[3,]	4	22	47	6	2307	51	9	170	23	595		
##	[4,]	2	23	12	1685	195	25	238	60	796	2		
##	[5,]	298	2167	1272	7	2601	5	516	988	43	8		
##	[6,]	1051	2	32	85	156	45	40	148	139	121		
##		[,129]	[,130]	[,131]	[,132]	[,133]	[,134]	[,135]	[,136]	[,137]	[,138]		
##	[1,]	4	226	65	16	38	1334	88	12	16	283		
##	[2,]	168	145	23	4	1690	15	16	4	1355	5		
##	[3,]	116	595	1352	13	191	79	638	89	2	14		
##	[4,]	4	671	7	2804	5	4	559	154	888	7		
##	[5,]	79	120	15	595	13	784	25	3171	18	165		
##	[6,]	664	665	10	10	1361	173	4	749	2	16		
##		[,139]	[,140]	[,141]	[,142]	[,143]	[,144]	[,145]	[,146]	[,147]	[,148]		
##	[1,]	5	16	4472	113	103	32	15	16	5345	19		

```
## [2,]      28      6      52     154     462      33      89      78     285      16
## [3,]       9      8     106     607     624      35     534       6     227       7
## [4,]    726     50      26      49    7008      15     566      30     579      21
## [5,]    170    143      19      14       5    7224       6     226     251       7
## [6,]   3804      8       4     226      65      12      43     127      24       2
##      [,149] [,150]
## [1,]    178     32
## [2,]    145     95
## [3,]    129    113
## [4,]     64   2574
## [5,]     61    113
## [6,]     10     10
```

```
# plot the distribution
hist(y_train, breaks = 2)
```



## Model Definition and Compilation

We define an RNN model with an embedding layer and LSTM layer:

```
model <- keras_model_sequential() %>%
  layer_embedding(input_dim = 10000, output_dim = 128) %>%
  layer_lstm(units = 64) %>%
  layer_dense(units = 1, activation = "sigmoid")

model %>% compile(
  loss = "binary_crossentropy",
  optimizer = "adam",
  metrics = c("accuracy")
)
```

## Model Training

We train the model for 20 epochs with a batch size of 100 samples.

```
model %>% fit(  
  x_train, y_train,  
  epochs = 20,  
  batch_size = 100,  
  validation_data = list(x_val, y_val)  
)
```

```
## Epoch 1/20  
## 1/1 - 6s - loss: 0.6931 - accuracy: 0.5100 - val_loss: 0.6929 - val_accuracy: 0.5149 - 6s/epoch - 6s  
## Epoch 2/20  
## 1/1 - 4s - loss: 0.6883 - accuracy: 0.6500 - val_loss: 0.6927 - val_accuracy: 0.5117 - 4s/epoch - 4s  
## Epoch 3/20  
## 1/1 - 4s - loss: 0.6831 - accuracy: 0.6800 - val_loss: 0.6926 - val_accuracy: 0.5062 - 4s/epoch - 4s  
## Epoch 4/20  
## 1/1 - 3s - loss: 0.6774 - accuracy: 0.6200 - val_loss: 0.6926 - val_accuracy: 0.5045 - 3s/epoch - 3s  
## Epoch 5/20  
## 1/1 - 4s - loss: 0.6708 - accuracy: 0.6000 - val_loss: 0.6927 - val_accuracy: 0.5036 - 4s/epoch - 4s  
## Epoch 6/20  
## 1/1 - 4s - loss: 0.6631 - accuracy: 0.6000 - val_loss: 0.6929 - val_accuracy: 0.5029 - 4s/epoch - 4s  
## Epoch 7/20  
## 1/1 - 4s - loss: 0.6539 - accuracy: 0.6000 - val_loss: 0.6933 - val_accuracy: 0.5028 - 4s/epoch - 4s  
## Epoch 8/20  
## 1/1 - 4s - loss: 0.6430 - accuracy: 0.6000 - val_loss: 0.6942 - val_accuracy: 0.5028 - 4s/epoch - 4s  
## Epoch 9/20  
## 1/1 - 3s - loss: 0.6300 - accuracy: 0.6000 - val_loss: 0.6960 - val_accuracy: 0.5028 - 3s/epoch - 3s  
## Epoch 10/20  
## 1/1 - 3s - loss: 0.6145 - accuracy: 0.6000 - val_loss: 0.6993 - val_accuracy: 0.5028 - 3s/epoch - 3s  
## Epoch 11/20  
## 1/1 - 3s - loss: 0.5961 - accuracy: 0.5900 - val_loss: 0.7058 - val_accuracy: 0.5028 - 3s/epoch - 3s  
## Epoch 12/20  
## 1/1 - 3s - loss: 0.5744 - accuracy: 0.5800 - val_loss: 0.7187 - val_accuracy: 0.5027 - 3s/epoch - 3s  
## Epoch 13/20  
## 1/1 - 3s - loss: 0.5499 - accuracy: 0.5800 - val_loss: 0.7463 - val_accuracy: 0.5027 - 3s/epoch - 3s  
## Epoch 14/20  
## 1/1 - 4s - loss: 0.5259 - accuracy: 0.5800 - val_loss: 0.7744 - val_accuracy: 0.5028 - 4s/epoch - 4s  
## Epoch 15/20  
## 1/1 - 3s - loss: 0.5014 - accuracy: 0.5800 - val_loss: 0.7609 - val_accuracy: 0.5028 - 3s/epoch - 3s  
## Epoch 16/20  
## 1/1 - 3s - loss: 0.4619 - accuracy: 0.6700 - val_loss: 0.7383 - val_accuracy: 0.5042 - 3s/epoch - 3s  
## Epoch 17/20  
## 1/1 - 4s - loss: 0.4229 - accuracy: 0.8100 - val_loss: 0.7254 - val_accuracy: 0.5085 - 4s/epoch - 4s  
## Epoch 18/20  
## 1/1 - 4s - loss: 0.3845 - accuracy: 0.8700 - val_loss: 0.7234 - val_accuracy: 0.5146 - 4s/epoch - 4s  
## Epoch 19/20  
## 1/1 - 3s - loss: 0.3387 - accuracy: 0.9200 - val_loss: 0.7433 - val_accuracy: 0.5220 - 3s/epoch - 3s  
## Epoch 20/20  
## 1/1 - 3s - loss: 0.2828 - accuracy: 0.9300 - val_loss: 0.9024 - val_accuracy: 0.5270 - 3s/epoch - 3s
```

## Model Evaluation

We evaluate the model on the test set:

```
loss_and_metrics <- model %>% evaluate(x_test, imdb$test$y, batch_size = 32)
```

```
## 782/782 - 9s - loss: 0.9021 - accuracy: 0.5250 - 9s/epoch - 12ms/step
```

```
cat("Test Loss:", loss_and_metrics[[1]], "\n")
```

```
## Test Loss: 0.9020851
```

```
cat("Test Accuracy:", loss_and_metrics[[2]], "\n")
```

```
## Test Accuracy: 0.525
```

## Conclusion

We have successfully trained and evaluated an RNN model for sentiment analysis on the IMDB dataset with the specified modifications. Further experiments can be conducted to explore the impact of using pre-trained word embeddings and varying the number of training samples on model performance.